

Deep Learning - Regularization

Itthi Chatnuntawech

Fully connected layer
(Dense)

Convolutional layer
Conv1D, 2D, 3D, ...
separable Conv

Optimizer
SGD
Adam
RMSprop

Pooling layer
max-pooling
average-pooling


Evaluation metric
accuracy
F1-score
AUC
confusion matrix

Deep Learning



Activation function
sigmoid
softmax
ESP (swish)
ReLU

Loss function
categorical crossentropy
binary crossentropy
mean squared error
mean absolute error

Regularization
Dropout
Data augmentation
 l_1, l_2 regularizations



- ❖ Combine basic components to build a neural network
 - More components → “More” representative power




```
# Import necessary modules
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
```

```
# Create the model
model = keras.Sequential()
model.add(layers.Dense(3, activation="relu"))
model.add(layers.Dense(2, activation="relu"))
model.add(layers.Dense(3, activation="softmax"))
```

```
# Compile the model
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.CategoricalCrossentropy(),
)
```



```
# Train the model for 100 epochs with a batch size of 32
model.fit(x_train, y_train, batch_size=32, epochs=100, validation_data=(x_val, y_val))
```



Model

Loss function
and optimizer

Regularization



- ❖ Regularization is frequently used to mitigate overfitting
 - ❖ Add a regularization term to the loss function
 - ❖ ℓ_2 regularization

$$\min \sum_{i=1}^N L(y_i, f_W(x)) \quad \rightarrow \quad \min \sum_{i=1}^N L(y_i, f_W(x)) + \lambda \|w\|_2^2$$


- ❖ ℓ_1 regularization

$$\min \sum_{i=1}^N L(y_i, f_W(x)) \quad \rightarrow \quad \min \sum_{i=1}^N L(y_i, f_W(x)) + \lambda \|w\|_1$$

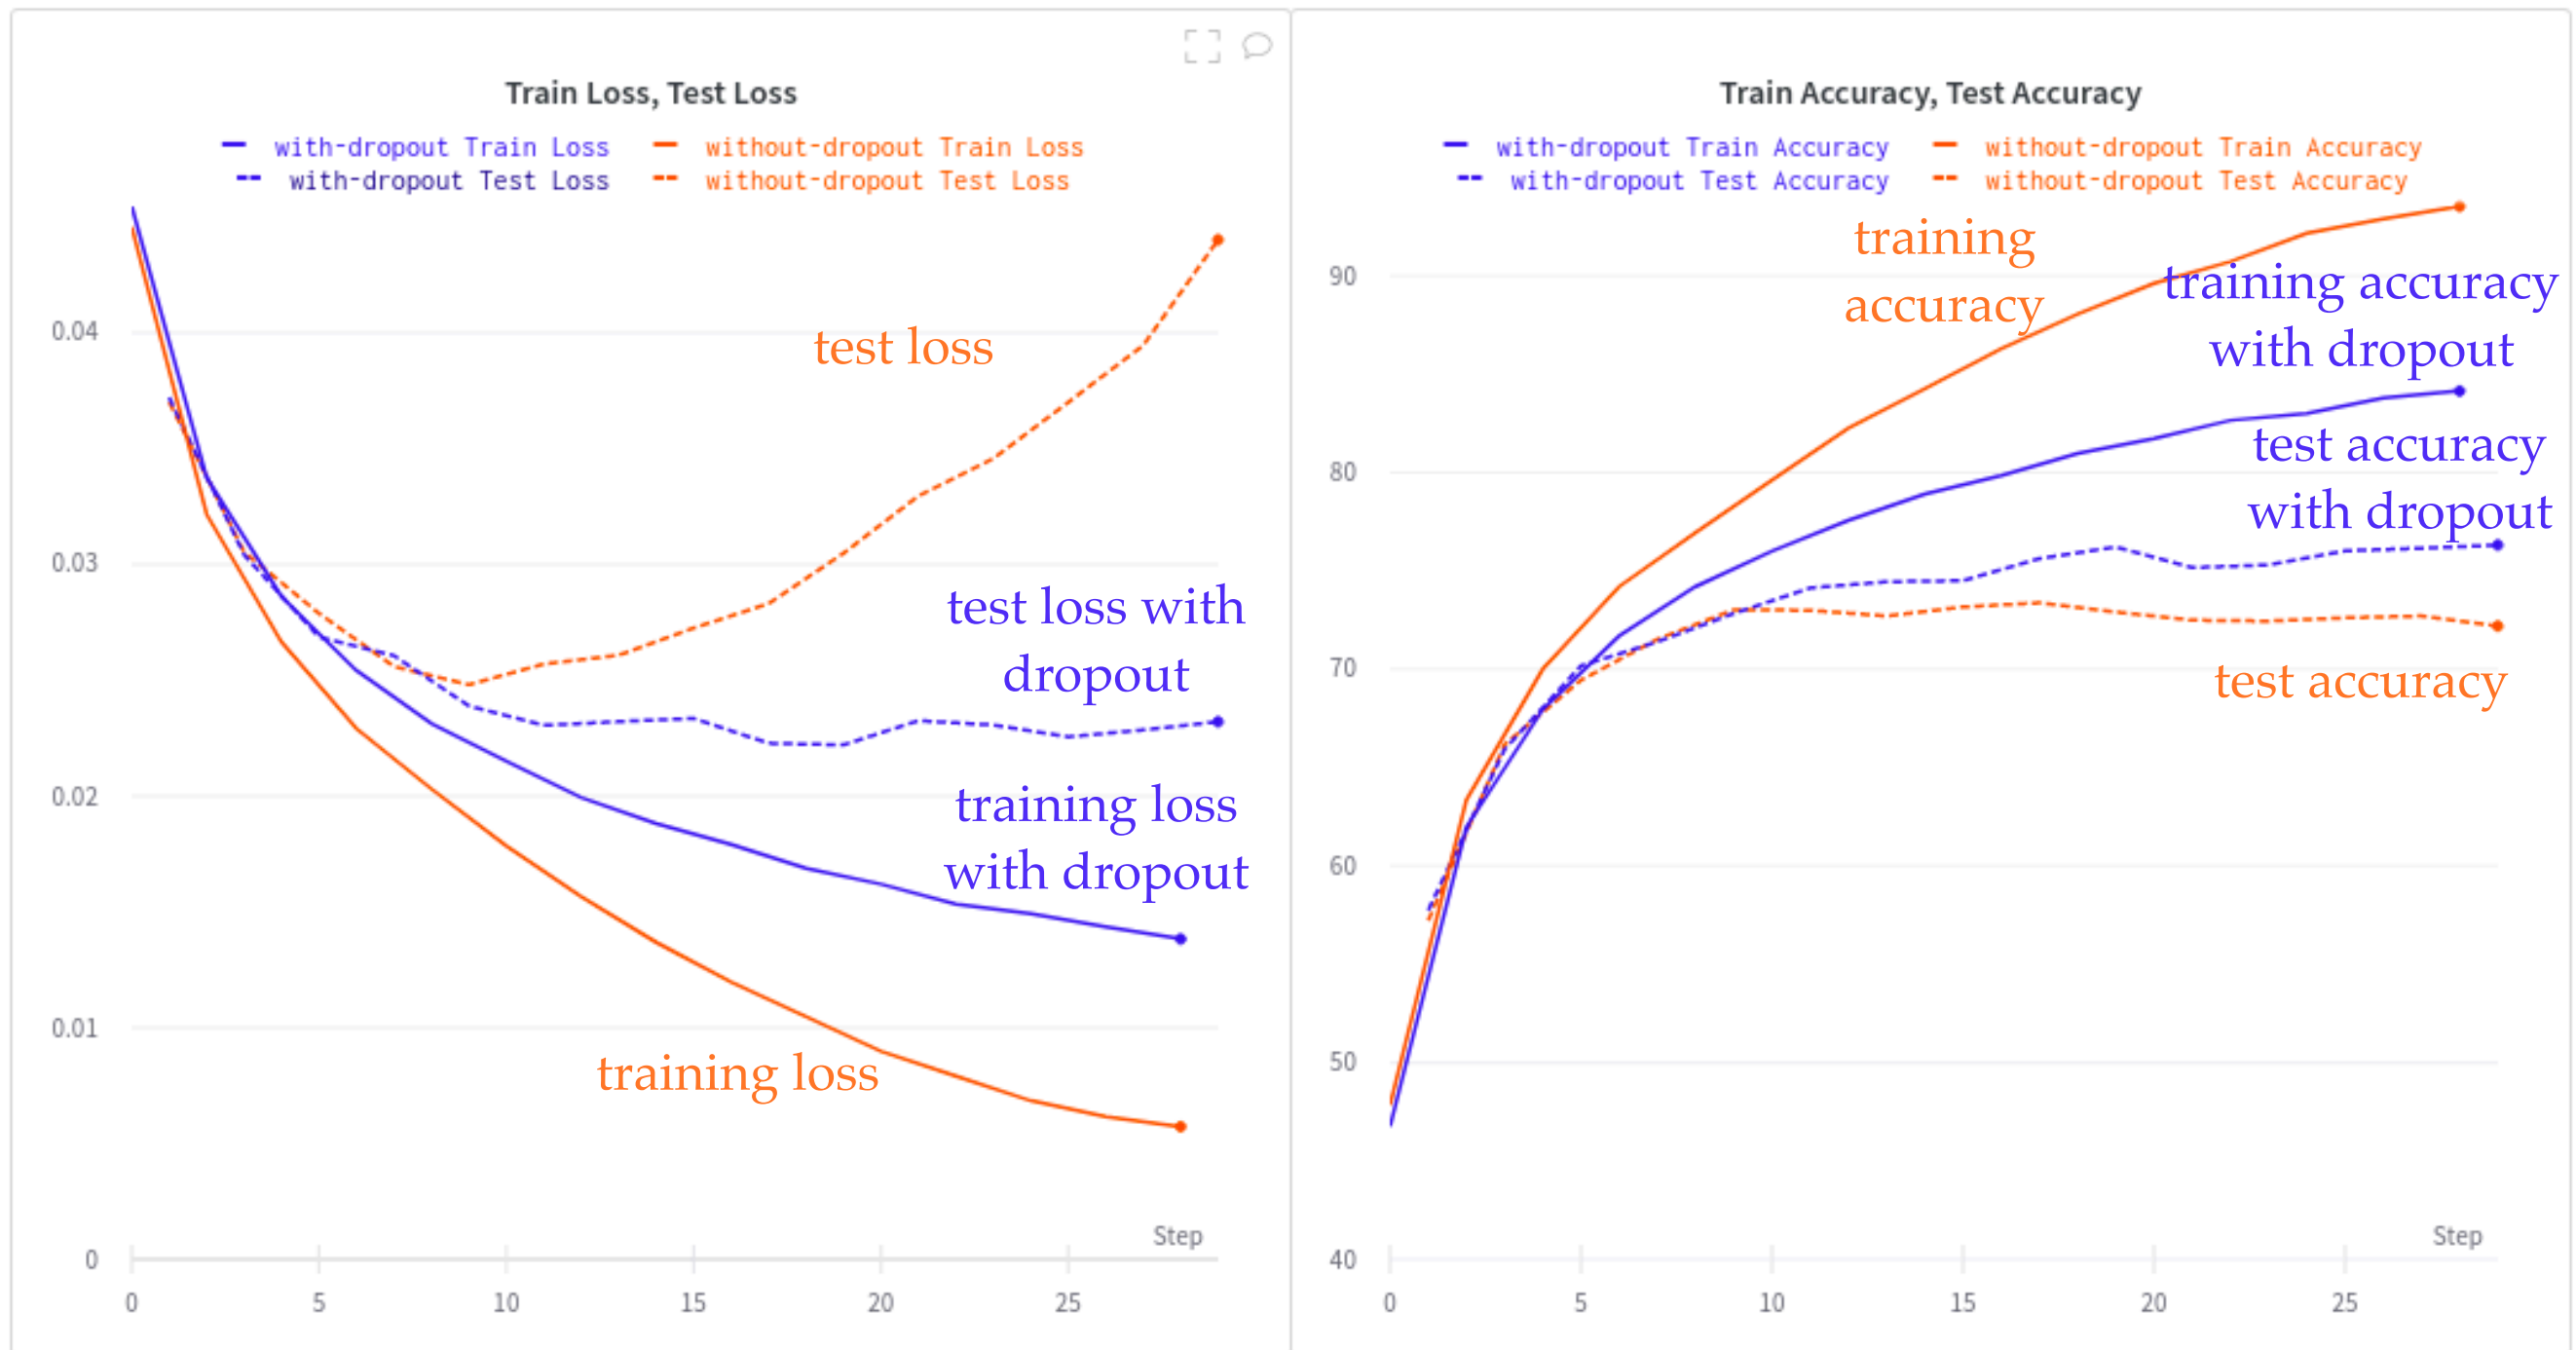
- ❖ Use dropout - much more popular

Dropout

- ❖ Dropout can be used to reduce overfitting
 - ❖ Randomly omit some neurons / units




Dropout




Batch Normalization

- ❖ Normalize every batch
 - ❖ reduce internal covariate shift problems*
- ❖ Can be thought of as a form of implicit regularization
 - ❖ can help reduce overfitting



Batch Normalization

- ❖ Normalize every batch
 - ❖ reduce internal covariate shift problems*
- ❖ Can be thought of as a form of implicit regularization
 - ❖ can help reduce overfitting



Batch Normalization — Speed up Neural Network Training

*This has been challenged by more recent work

Batch Normalization

- ❖ Reduce internal covariate shift problems*

batch size # of features

Input: $x : N \times D$

Learnable params:

$\gamma, \beta : D$

Intermediates: $\mu, \sigma : D$
 $\hat{x} : N \times D$

Output: $y : N \times D$

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

- ❖ More robust to bad initialization

Data Augmentation

